

Predicting risk of drug consumption based on personality features, demographic variables and previous consumption.

Camila Batista Rahal and Elena Maestre Soteras

Abstract

Drug consumption and misuse remain significant public health concerns, shaped by an interaction of personality traits, demographic characteristics, and historical behavior. The study employs a variety of predictive models (logistic regression, lasso, random forest, and Xgboost) to analyze the relationships between this personality factors, demographic variables, and drug consumption patterns, aiming to identify significant predictors for the risk of future drug use. The participants were classified as 'consumer' or 'non-consumer.' XGBoost baseline achieve the best performance with a recall of 84.3%, followed by Random Forest with recall of 83%. Random Forest achieves the highest AUC (0.90), making it the most effective model for distinguishing between consumers and non-consumers classes. Bagging Classifier and Optimized Random Forest closely follow with an AUC of 0.89, indicating strong performance with balanced sensitivity and specificity. This research demonstrates the potential of data-driven methods for understanding and addressing drug use behaviors, while offering a foundation for future studies that integrate psychological and demographic data.

Exploratory Data Analysis

The dataset had 1,885 participants, no missing data, and the features were quantified and standardized. The dataset shows a high prevalence of participants aged 18–24, a predominance of consumers within this group, a disproportionate number of White participants compared to other ethnicities, and a higher number of male consumers than female (Table 1). Summary statistics (Table 2) confirm key predictors like age, country, ethnicity, impulsivity, and sensation-seeking. Lower education appears linked to higher use. Target imbalance (mean = 0.552) was handled with SMOTE. Figure 1 below visualizes the distribution of drug usage frequencies across various substances. It highlights trends such as a large number of participants reporting never having used heroin or crack, while substances like ecstasy and mushrooms show relatively higher recent use. Additionally, there is a high prevalence of users for substances like chocolate, nicotine, caffeine, alcohol, and cannabis. These substances were excluded when creating the target variable.

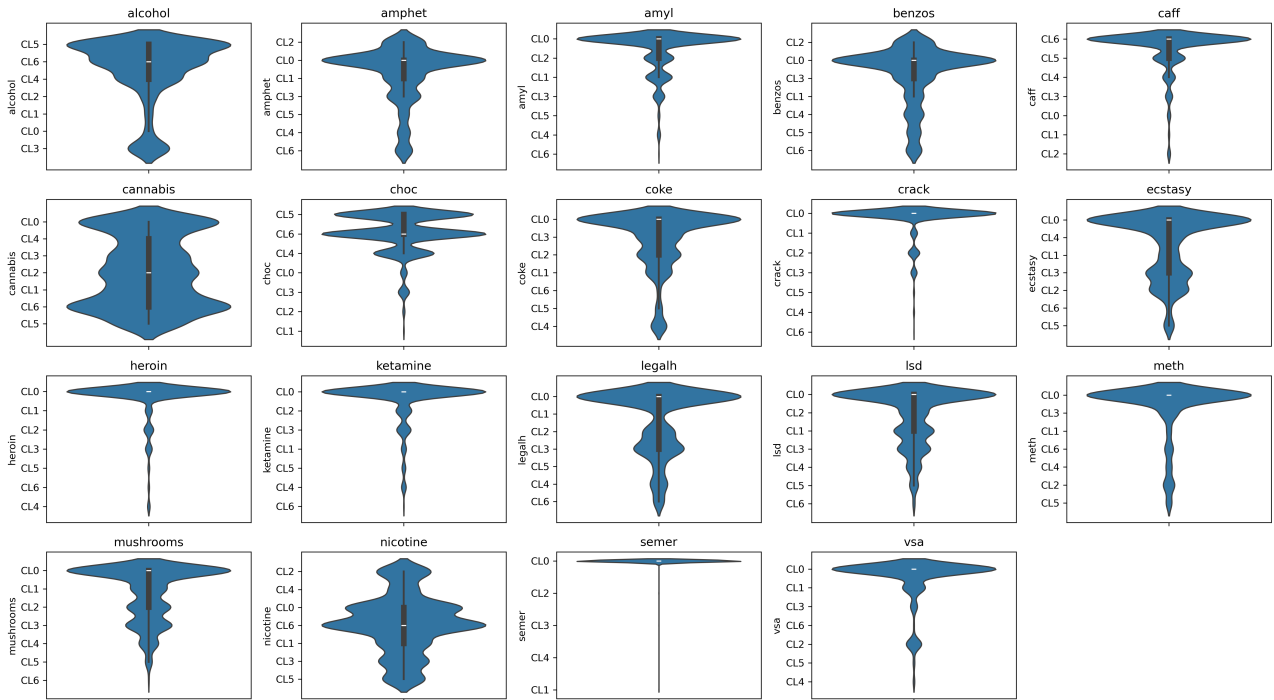


Figure 1: Violin plot of all drugs consuming behavior from CL0 (Never used) to CL6 (Used last day).

To analyze the correlation between features a correlation matrix was performed (Figure 2). It demonstrated that "Country" and "Age" show strong negative correlations with the target (-0.55 and -0.42), while "Sensation-Seeking" (ss) and "Impulsivity" have positive correlations (0.44 and 0.32), indicating their influence on user

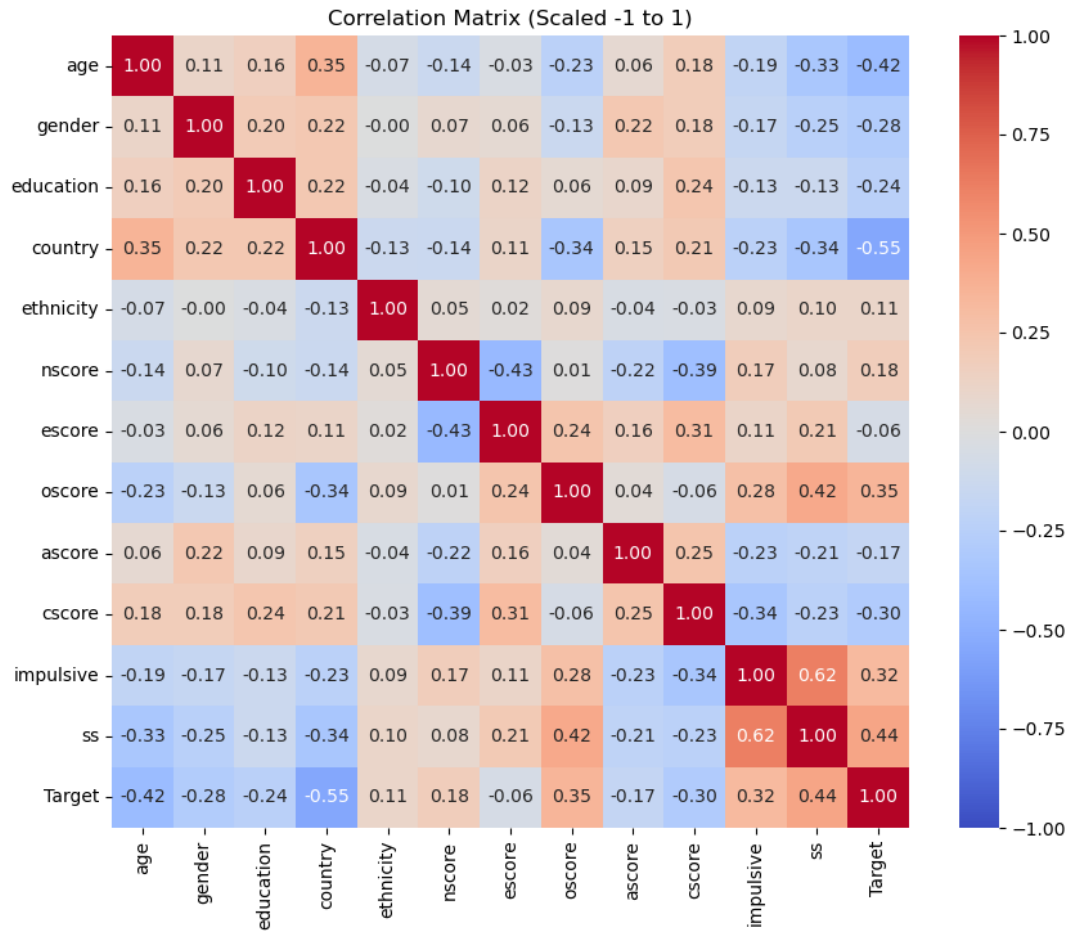


Figure 2: Correlation matrix of features.

classification. Additionally, "nscore" and "escore" are moderately negatively correlated (-0.43), and "ss" and "Impulsive" are strongly positively correlated (0.62), suggesting feature interdependencies. This analysis hints the most important features it can be used as a predictor for the model.

Data Preparation

- Data used was preprocessed previously.
- 5% of the data was reserved as final evaluation set *Hold-Out, and the remaining data was split into 80% training, 10% validation, 10% testing. Random seeds ensured reproducibility.
- Target was binary (consumer vs non-consumer) and 11 features: the Big Five personality traits (NEO-FFI-R), those are nscore (Neutoticism), escore (Extraversion), oscore (Openness to Experience), ascore (Agreeableness), and cscore (Conscientiousness). Impulsivity (BIS-11), sensation seeking (ImpSS), and demographic information: age, gender, education, and country.
- Data was standardized.

Models implemented

- Logistic Regression with L1 regularization
- Random Forest (baseline)
- Random Forest (Optimized with hyperparameters)
- XGBoost (baseline)

- XGBoost (Optimized with hyperparameters)
- SVM
- Decision Tree
- Bagging Classifier

Model evaluation metrics

Each models performance was evaluated using, besides the validation and hold-out set, accuracy, precision and recall, f1 score, cross-validation *only on 3 models, and ROC-AUC. In this project, which focuses on predicting the risk of drug consumption based on personality features, demographic variables, and previous consumption, the primary evaluation metric is recall. Recall is prioritized because the goal is to ensure that individuals who are actual consumers are accurately identified, minimizing false negatives. This is particularly important as missing potential consumers (false negatives) could lead to missed opportunities for intervention. While recall is the primary focus, precision is also monitored to ensure that the model maintains reasonable targeting, avoiding an excessive number of false positives. Additionally, F1-Score is used as a balanced metric to evaluate the trade-off between precision and recall, and ROC-AUC is included to assess the overall discriminatory power of the models.

Feature importance analysis

- **Logistic Regression with L1 regularization:** "ss" has the highest positive coefficient, while "age" has a significant negative coefficient, indicating their contrasting roles in prediction. Lasso excluded impulsiveness.
- **Linear SVM:** "ss" is the most significant positive predictor, "age" contributes negatively, others are smaller.
- **Bagging Classifier:** "Age" and "ss" dominate, age is strongest.
- **Decision Tree:** "Age" is dominant, followed by "ss".
- **Random Forest:** Both "ss" and "age" are key predictors.
- **Optimized XGBoost:** "Age" is top, followed by "ss".
- **Optimized Random Forest:** "ss" and "age" show balanced importance.
- **Conclusion:** "ss" and "age" are most important across models.

Results

Model performance

When focusing on recall, the Baseline XGBoost achieves the highest recall (84.3%) followed by baseline Random Forest on the validation set, and the Bagging Classifier and Baseline XGBoost show the highest recall (80.7%) on the hold-out set (Table 3). These results suggest that these models are particularly effective in identifying positive cases (consumers). The inconsistency of the results between the validation set and the hold-out, and the lower performance of the optimized models against the baseline models shows that there still space for improvement.

Cross-validation

The evaluation used Stratified K-Fold Cross-Validation (5 folds) for Logistic Regression and during manual hyperparameter tuning for XGBoost, while GridSearchCV with 5-fold cross-validation was employed for the optimized Random Forest and the optimized XGBoost.

Cross-Validation metrics for Logistic Regression with Lasso: Mean Precision: 0.8128 Mean Recall: 0.7877 Mean F1 Score: 0.8000 Mean ROC-AUC: 0.8670

Cross-Validation metrics for Random Forest (Optimized): Best cross-validation F1 score: 0.8062

Cross-Validation for XGBoost (Optimized): Mean Logloss: 0.4471 ± 0.0276 Mean AUC: 0.8762 ± 0.0139

The result shows that XGBoost performs well, with accurate probability predictions (logloss: 0.4471) and strong class distinction (AUC: 0.8762), with consistent performance across folds (low standard deviations).

ROC AUC

Random Forest achieved the highest AUC (0.90), followed by Bagging and Optimized Random Forest (0.89). Optimized XGBoost (0.88) outperformed its baseline (0.86). Logistic Regression and SVM showed solid results (0.86–0.87), while Decision Tree had the lowest AUC (0.85).

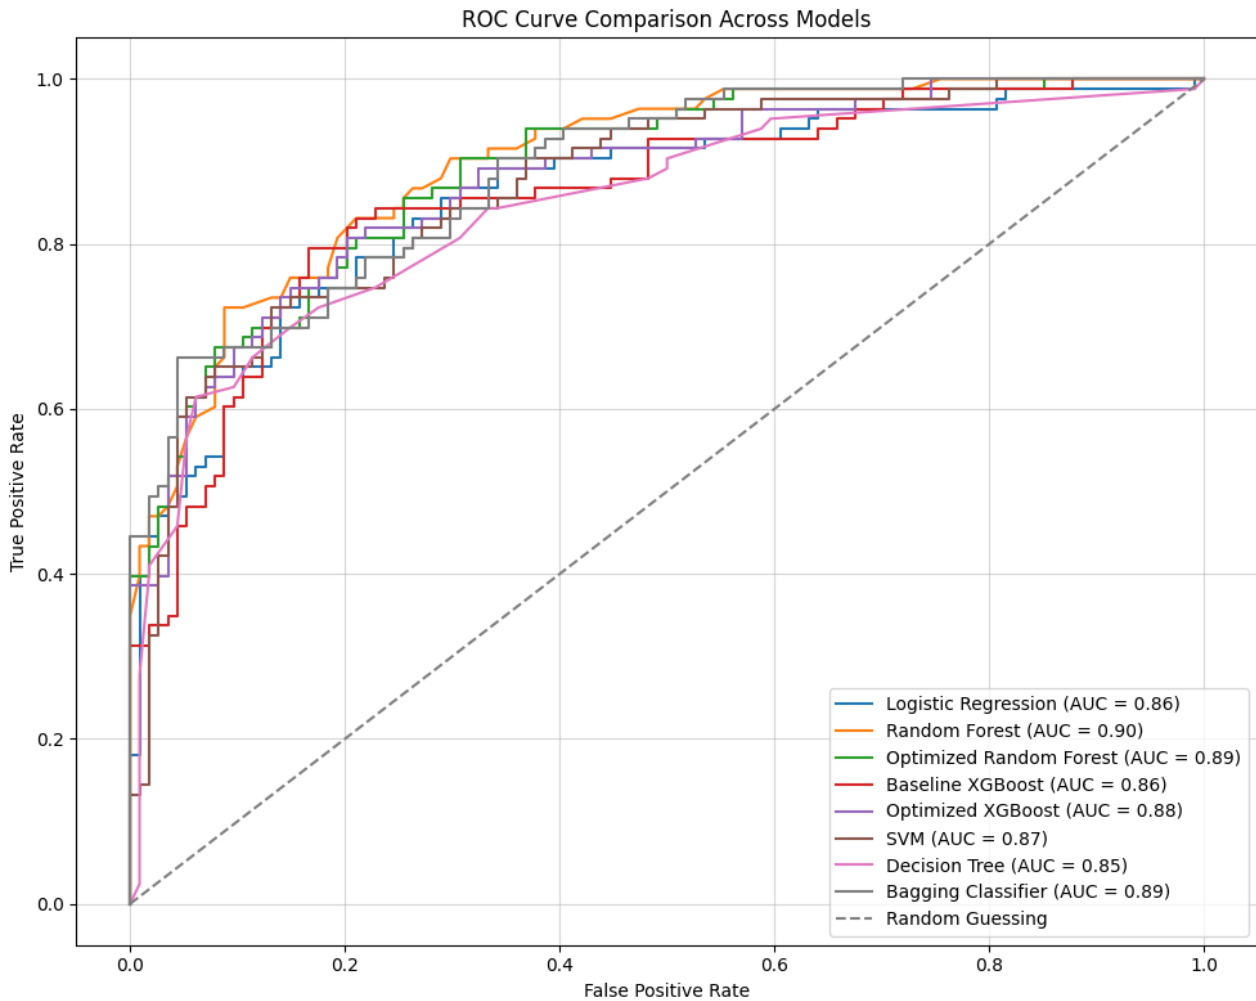


Figure 3: ROC Curve comparison across supervised machine learning models.

Discussion

Model performance consistency suggests strong data patterns and stable predictors. Sensation seeking (ss) and age consistently emerged as the top predictors across models. Logistic Regression and SVM highlighted ss positively and age negatively. Tree-based models also emphasized these features, with Optimized XGBoost and Random Forest showing strong generalization. Despite existing evidence, Lasso excluded impulsivity, possibly due to multicollinearity or weak individual effect.

Implications

Results support using psychological and demographic data for early risk identification in public health strategies.

Limitations

The datasets imbalance and youth-skewed sample may limit generalizability. Broader, more diverse data is needed in future research. Excluding variables may have overlooked interactions between legal and illegal substances. The static nature of the data also omits temporal trends that could inform targeted interventions. Future studies should include older, more diverse populations and explore dynamic usage patterns for improved insight and model performance.

Conclusion

This study demonstrates that personality traits—especially sensation seeking—and age are key predictors of psychoactive drug use, with models like XGBoost and Random Forest achieving strong performance. Integrating psychological and demographic data into predictive models offers valuable tools for early identification and targeted public health interventions.

Tables

Table 1: Demographics Summary

Target	Non-consumer	Consumer
Age category		
18-24 years	113	524
25-34 years	218	262
35-44 years	215	140
45-54 years	211	83
55-64 years	70	23
65+ years	13	5
Gender category		
Female	552	385
Male	288	652
Education category		
Doctorate degree	53	36
Left school at 16 years	53	45
Left school at 17 years	11	18
Left school at 18 years	25	74
Left school before 16 years	14	14
Masters degree	189	94
Professional certificate/diploma	149	121
Some college or university (no certificate)	85	418
University degree	261	217
Country category		
Australia	8	44
Canada	32	55
New Zealand	1	4
Other	28	90
Republic of Ireland	6	14
UK	716	328
USA	49	502
Ethnicity category		
Asian	17	8
Black	25	8
Mixed-Black/Asian	0	3
Mixed-White/Asian	8	12
Mixed-White/Black	7	12
Other	19	43
White	764	951

Table 2: Summary statistics of features and target variable

	count	mean	std	min	25%	50%	75%	max
Age	1877	0.037	0.878	-0.951	-0.951	-0.078	0.497	2.591
Gender	1877	-0.0007	0.482	-0.482	-0.482	-0.482	0.482	0.482
Education	1877	-	0.949	-2.435	-0.611	-0.059	0.454	1.984
Country	1877	0.358	0.699	-0.57	-0.57	0.96	0.96	0.96
Ethnicity	1877	-0.309	0.166	-1.107	-0.316	-0.316	-0.316	1.907
Nscore	1877	0.00055	0.998	-3.464	-0.678	0.042	0.629	3.273
Escore	1877	-0.0019	0.997	-3.273	-0.695	0.003	0.637	3.273
Oscore	1877	-0.0032	0.995	-3.273	-0.717	-0.019	0.723	2.901
Ascore	1877	0.00065	0.996	-3.464	-0.606	-0.017	0.76	3.464
Cscore	1877	0.00039	0.997	-3.464	-0.652	-0.006	0.584	3.464
Impulsive	1877	0.0052	0.954	-2.555	-0.711	-0.217	0.529	2.901
Ss	1877	-0.0074	0.962	-2.078	-0.525	0.079	0.765	1.921
Target	1877	0.552	0.497	0	0	1	1	1

Table 3: Model performance across datasets

Model	Accuracy	Precision	Recall	F1 Score	AUC	Dataset
Logistic Regression	0.771574	0.706522	0.783133	0.742857	0.859121	Validation
Random Forest	0.791878	0.718750	0.831325	0.770950	0.898594	Validation
Optimized Random Forest	0.791878	0.733333	0.795181	0.763006	0.889347	Validation
Baseline XGBoost	0.802030	0.729167	0.843373	0.782123	0.858592	Validation
Optimized XGBoost	0.802030	0.744444	0.807229	0.774566	0.876242	Validation
SVM	0.766497	0.690722	0.807229	0.744444	0.869267	Validation
Decision Tree	0.761421	0.704545	0.746988	0.725146	0.846121	Validation
Bagging Classifier	0.761421	0.691489	0.783133	0.734463	0.887022	Validation
Logistic Regression	0.682692	0.679245	0.692308	0.685714	0.793639	Hold-Out
Random Forest	0.798077	0.816327	0.769231	0.792079	0.891272	Hold-Out
Optimized Random Forest	0.807692	0.820000	0.788462	0.803922	0.876849	Hold-Out
Baseline XGBoost	0.817308	0.823529	0.807692	0.815534	0.884985	Hold-Out
Optimized XGBoost	0.807692	0.820000	0.788462	0.803922	0.876898	Hold-Out
SVM	0.759615	0.764706	0.750000	0.757282	0.837278	Hold-Out
Decision Tree	0.769231	0.791667	0.730769	0.760000	0.846524	Hold-Out
Bagging Classifier	0.788462	0.777778	0.807692	0.792453	0.866124	Hold-Out